

Docker quickguide

Data Engineering Frühlingssemester 2020

Marcel Huber

Version 20.4.0, 2020-04-06: updated Datastructures and Indexes

Inhaltsverzeichnis

Docker	2
Installation.....	2
Special note on how to use X client containers (linux, mac)	3
Basic commands and images.....	4

This document describes some of the helper tools used within this course.

Docker

Instead of installing the tools as native packages for your operating system, we can make use of [Docker](#) images for the different services we use.

To get a short introduction into Docker including some basic examples, check out [these slides](#) of the [Cloud Solutions](#) lecture.

Installation

Linux

Check the [Linux installation page](#) on how to install docker community edition after having checked the [prerequisites](#).

Mac

Check the [Mac installation page](#) on how to install docker community edition after having checked the prerequisites.



Docker for Mac requires OS X El Capitan 10.11 or newer macOS release running on a 2010 or newer Mac, with Intel's hardware support for MMU virtualization. The app will run on 10.10.3 Yosemite, but with limited support. Please see [What to know before you install](#) for a full explanation and list of prerequisites.



If you want to [share files outside the /Users directory](#) with your container, you have to add these directories in the docker ->**Preferences**->**File sharing**.

Windows

Check the [Windows installation page](#) on how to install docker community edition after having checked the prerequisites.



Docker for Windows requires 64bit Windows 10 Pro and Microsoft Hyper-V. Please see [What to know before you install](#) for a full list of prerequisites.



- If you want to [share directories](#) with your container, you have to add these directories in the docker ->**Settings**->**Shared Drives**.
- If you get errors starting your container when using **--publish** run option, you might need to disable **Experimental Features** in the docker ->**Settings** ->**Daemon**. ([this issue](#))

Special note on how to use X client containers (linux, mac)

In order to access your X server from within a docker X client, follow these steps based on your platform.

Linux

1. Pass your **DISPLAY** variable when creating/running the container

```
--env DISPLAY=unix$DISPLAY
```

2. Map your X11 socket as volume into the container

```
--volume /tmp/.X11-unix:/tmp/.X11-unix
```

3. Allow access to your X11 server from local clients

```
# allow local clients only
xhost local:
# or allow all
xhost +
```

Mac

According to [this issue comment](#), the following steps are needed.

1. Start **socat** to expose local xquartz socket on a TCP port

```
socat TCP-LISTEN:6000,reuseaddr,fork UNIX-CLIENT:\"$DISPLAY\"
```

This is insecure on public networks, add **bind**, **su** and **range** options to **socat** to limit access.

2. Pass the display to container (assuming virtualbox host is available on 192.168.59.3):

```
--env DISPLAY=192.168.59.3:0
```

Windows (unverified!)

The following procedure, as stated [here](#) might theoretically work but isn't verified yet.

1. Install [Xming](#)
2. Start Xming like this, **-ac** disables access control checks:

```
Xming -ac -multiwindow -clipboard
```

3. Pass the display to the container:

```
--env DISPLAY=192.168.59.3:0.0
```

Replace the above IP address with your hosts current address.

Basic commands and images

show available images

```
docker image ls
```

show running and dead containers

```
docker container ls -a
```

PostgreSQL + PostGIS + cstore_fdw

This image contains the latest postgres server including the cli tool **psql**. Additionally, postgis and **cstore_fdw** extensions are installed and ready to use.

Build image

```
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory Helpers/
  build postgres
```

Run the server in the background and make it accessible from the host at localhost:5432:

```
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory .
  up -d
  postgres
```

① -d runs the service in the background

Use the `psql` cli tool to access the server

```
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory . ①
run postgres ②
psql ③
  -h postgres ④
  -U postgres ⑤
```

- ① Bind mount the specified directory to `/src` in the container.
- ② Run service `postgres` interactively in the foreground.
- ③ Instead of the default command, execute `psql ...`.
- ④ Hostname of the `postgres` database.
- ⑤ Username of the dba, usually `postgres`.

Execute (p)sql commands from a file which resides in your local file system.

To access files from the local file system, we need to map it using the `--volume` option. The service entry in the compose file already maps the `project-directory` as `/src` into the container.

```
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory .
run postgres
psql
  -h postgres
  -U postgres
  -v ON_ERROR_STOP=1 ①
  -f /src/Databases/bank/0_runAllScripts.sql ②
```

- ① Abort the script in case of errors. Default is to silently continue running the script which might lead to a non-working database in case of errors.
- ② Full path name to the SQL script file to import. This path needs to be present within the container.

Eclipse-Java-EE

To update or manually build an eclipse-jee image use the following command.

Build an updated image

```
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory Helpers/
  build eclipse-jee
```

Run eclipse in the background

```
xhost local:
docker-compose
  --file Helpers/docker-compose.yml
  --project-name dataeng
  --project-directory .
  up -d eclipse-jee
```

If you place your workspace directory below the default folder */nobody*, changes will automatically be persisted between sessions. The corresponding volume is probably named *dataeng_eclipse-jee-data*. Plugins and updates to them will also be stored in that volume when updated.

The following tools will help you to keep your code clean but are not pre-installed. You can easily find and install them using Eclipse-Marketplace.

Useful tools

- EclEmma
- Checkstyle Plug-in
- PMD Plug-in
- FindBugs Feature
- LiClipseText Feature